

```
import pandas as pd
import numpy as np
import os
import joblib
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, roc_auc_score, classification_report,
    precision_recall_curve
)
```

```
from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from xgboost import XGBClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.calibration import CalibratedClassifierCV
```

1. Load dataset

```
df = pd.read_csv("/content/balanced_crime_data.csv")
```

2. Define target + drop irrelevant cols

```
target = "Arrest"
```

```
drop_cols = ["ID", "Case Number", "Date", "Updated On", "Location"]
df = df.drop(columns=[col for col in drop_cols if col in df.columns])
```

3. Split features & target

```
X = df.drop(columns=[target])
y = df[target].astype(int)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

4. Preprocessor

```
numeric_features = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_features = X.select_dtypes(exclude=[np.number]).columns.tolist()
```

```
numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])
```

```
categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore"))
])
```

```
preprocessor = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)
```

5. Base models with balancing

```
rf = RandomForestClassifier(
    n_estimators=200, random_state=42, n_jobs=-1,
    class_weight="balanced"
)
```

```
xgb = XGBClassifier(
    n_estimators=300, learning_rate=0.05, random_state=42,
    n_jobs=-1, use_label_encoder=False, eval_metric="logloss",
    scale_pos_weight=(len(y_train) - sum(y_train)) / sum(y_train))
```

)

```
base_models = [("rf", rf), ("xgb", xgb)]
```

6. Meta model (MLP) with calibration

```
meta_model = MLPClassifier(  
    hidden_layer_sizes=(64, 32),  
    activation="relu",  
    solver="adam",  
    max_iter=300,  
    random_state=42,  
    early_stopping=True
```

)

```
calibrated_meta = CalibratedClassifierCV(meta_model, method="isotonic", cv=3)
```

7. Stacking classifier

```
stacking_clf = StackingClassifier(  
    estimators=base_models,  
    final_estimator=calibrated_meta,  
    passthrough=True,  
    n_jobs=-1
```

)

8. Full pipeline

```
pipeline = Pipeline(steps=[  
    ("preprocessor", preprocessor),  
    ("stacking", stacking_clf)  
])
```

9. Train

```
print("Training stacking classifier...")  
pipeline.fit(X_train, y_train)
```

```

# -----
# 10. Create directory & save model
# -----

model_dir = "/mnt/data/saved_models"
model_filename = "crime_stacking_model_v2.pkl"
os.makedirs(model_dir, exist_ok=True)
model_path = os.path.join(model_dir, model_filename)

joblib.dump(pipeline, model_path)
print(f"Model directory ready: {model_dir}")
print(f"Model successfully saved at: {model_path}")

# 11. Evaluate
y_pred = pipeline.predict(X_test)
y_proba = pipeline.predict_proba(X_test)[:, 1]

print("\nClassification Report (Default Threshold 0.5):\n", classification_report(y_test,
y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
print(" Precision:", precision_score(y_test, y_pred))
print(" Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print(" ROC-AUC:", roc_auc_score(y_test, y_proba))

# 12. Threshold tuning function

def tune_threshold(y_true, y_proba, metric="f1"):
    precisions, recalls, thresholds = precision_recall_curve(y_true, y_proba)
    best_threshold, best_score = 0.5, 0
    for p, r, t in zip(precisions, recalls, thresholds):
        if metric == "f1":
            score = 2 * (p * r) / (p + r + 1e-6)
        elif metric == "recall":
            score = r
        elif metric == "precision":
            score = p

```

```
else:
    raise ValueError("Metric must be 'f1', 'recall', or 'precision'")
if score > best_score:
    best_threshold, best_score = t, score
return best_threshold, best_score
```

13. Tune threshold for F1

```
best_t, best_f1 = tune_threshold(y_test, y_proba, metric="f1")
print(f"\n Best Threshold for F1: {best_t:.3f} | F1 Score: {best_f1:.4f}")
```

Apply tuned threshold

```
y_pred_tuned = (y_proba >= best_t).astype(int)
print("\n Classification Report (Tuned Threshold):\n", classification_report(y_test,
y_pred_tuned))
```

14. Optional: plot precision–recall vs threshold

```
precisions, recalls, thresholds = precision_recall_curve(y_test, y_proba)
plt.figure(figsize=(8,6))
plt.plot(thresholds, precisions[:-1], label="Precision")
plt.plot(thresholds, recalls[:-1], label="Recall")
plt.axvline(x=best_t, color="red", linestyle="--", label=f"Best Threshold {best_t:.2f}")
plt.xlabel("Threshold")
plt.ylabel("Score")
plt.title("Precision-Recall vs Threshold")
plt.legend()
plt.grid(True)
plt.show()
```